

Practical Machine Learning

Amir Yazid

Thursday, March 17, 2016

Data preprocessing

```
library(caret)

## Loading required package: lattice

## Loading required package: ggplot2

library(rpart)
library(rpart.plot)
library(randomForest)

## randomForest 4.6-12

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:ggplot2':
##
##   margin

library(corrplot)
```

Read Data

```
trainRaw <- read.csv("D:/coursera/data/8 Practical Machine Learning/pml-training.csv")
testRaw <- read.csv("D:/coursera/data/8 Practical Machine Learning/pml-testing.csv")

dim(trainRaw)

## [1] 19622 160

dim(testRaw)

## [1] 20 160
```

Data Cleaning

```
##remove NA values
```

```
trainRaw <- trainRaw[, colSums(is.na(trainRaw)) == 0]
testRaw <- testRaw[, colSums(is.na(testRaw)) == 0]
```

##get rid of some columns that do not contribute much to the accelerometer measurements

```
classe <- trainRaw$classe
trainRemove <- grepl("^X|timestamp|window", names(trainRaw))
trainRaw <- trainRaw[, !trainRemove]
trainCleaned <- trainRaw[, sapply(trainRaw, is.numeric)]
trainCleaned$classe <- classe
testRemove <- grepl("^X|timestamp|window", names(testRaw))
testRaw <- testRaw[, !testRemove]
testCleaned <- testRaw[, sapply(testRaw, is.numeric)]
```

##split the cleaned training set into a pure training data set (70%) and a validation data set (30%)
##use the validation data set to conduct cross validation in future steps

```
set.seed(22519) # For reproducible purpose
inTrain <- createDataPartition(trainCleaned$classe, p=0.70, list=F)
trainData <- trainCleaned[inTrain, ]
testData <- trainCleaned[-inTrain, ]
```

Data Modelling

##Fit a predictive model for activity recognition using Random Forest algorithm
##because it automatically selects important variables and is robust to correlated covariates & outliers
##We will use 5-fold cross validation when applying the algorithm.

```
controlRf <- trainControl(method="cv", 5)
modelRf <- train(classe ~ ., data=trainData, method="rf", trControl=controlRf, ntree=250)
modelRf
```

```
## Random Forest
##
## 13737 samples
##    52 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 10989, 10989, 10991, 10990, 10989
## Resampling results across tuning parameters:
##
##  mtry  Accuracy   Kappa      Accuracy SD   Kappa SD
##    2    0.9908278 0.9883961 0.001675093   0.002121380
##   27    0.9911190 0.9887646 0.001755971   0.002222771
##   52    0.9840572 0.9798290 0.003497420   0.004425063
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 27.
```

```
##Then estimate the performance of the model on the validation data set
```

```
predictRf <- predict(modelRf, testData)
confusionMatrix(testData$classe, predictRf)
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction    A     B     C     D     E
##           A 1672     1     0     0     1
##           B     5 1129     5     0     0
##           C     0     1 1020     5     0
##           D     0     0    13   949     2
##           E     0     0     1     6 1075
```

```
## Overall Statistics
```

```
##
##           Accuracy : 0.9932
##           95% CI : (0.9908, 0.9951)
##           No Information Rate : 0.285
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9914
##           McNemar's Test P-Value : NA
```

```
##
## Statistics by Class:
```

```
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9970   0.9982   0.9817   0.9885   0.9972
## Specificity      0.9995   0.9979   0.9988   0.9970   0.9985
## Pos Pred Value    0.9988   0.9912   0.9942   0.9844   0.9935
## Neg Pred Value    0.9988   0.9996   0.9961   0.9978   0.9994
## Prevalence       0.2850   0.1922   0.1766   0.1631   0.1832
## Detection Rate    0.2841   0.1918   0.1733   0.1613   0.1827
## Detection Prevalence 0.2845   0.1935   0.1743   0.1638   0.1839
## Balanced Accuracy 0.9983   0.9981   0.9902   0.9927   0.9979
```

```
##calculate accuracy
```

```
accuracy <- postResample(predictRf, testData$classe)
accuracy
```

```
## Accuracy      Kappa
## 0.9932031 0.9914024
```

```
##calculate estimated out-of-sample error
```

```
oose <- 1 - as.numeric(confusionMatrix(testData$classe, predictRf)$overall[1])
oose
```

```
## [1] 0.006796941
```


